

# Zip-zip Trees: Making Zip Trees More Balanced, Biased, Compact, or Persistent

Ofek Gila<sup>1</sup> , Michael T. Goodrich<sup>1</sup> , and Robert E. Tarjan<sup>2</sup>

<sup>1</sup>University of California, Irvine

<sup>2</sup>Princeton University

WADS, 2023

# Motivation

- 2018 – Zip tree (Tarjan, Levy, Timmel)
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node

# Motivation

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node

# Motivation

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement

- 2018 – Zip tree (Tarjan, Levy, Timmel) - **asymmetric**
  - **1.5**  $\log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement
  - ✓ Strongly history independent

# Motivation

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement
  - ✓ Strongly history independent
  - ✓ May be partially persistent



# Motivation

- 2018 – Zip tree (Tarjan, Levy, Timmel) - asymmetric
  - $1.5 \log n$  expected average depth
    - Min key:  $0.5 \log n$ , Max key:  $\log n$
  - Space cost:  $\log \log n$  bits per node
- 2023 – Zip-zip trees:
  - ✓  $1.39 \log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement
  - ✓ Strongly history independent
  - ✓ May be partially persistent
  - ✓ Supports biased keys (still  $O(\log \log n)$  bits per node)

# Skip List




Figure 1: A sorted linked list

# Skip List

- What if you add 'fast lanes'?




Figure 1: A sorted linked list

# Skip List

- What if you add 'fast lanes'?




Figure 1: A skip list with one coin flip

# Skip List

- What if you add 'fast lanes'?




Figure 1: A skip list with two coin flips

# Skip List

- What if you add 'fast lanes'?




Figure 1: A skip list with three coin flips

# Skip List

- What if you add 'fast lanes'?




Figure 1: A skip list with four coin flips

# Skip List

- What if you add 'fast lanes'?
- Idea: 1 move in level  $k \approx 2$  in level  $k - 1 \approx 2^k$  in level 0




Figure 1: A skip list with four coin flips



# Skip List

- What if you add 'fast lanes'?
- Idea: 1 move in level  $k \approx 2$  in level  $k - 1 \approx 2^k$  in level 0
- $\mathcal{O}(\log n)$  expected search time




Figure 1: A skip list with four coin flips

# Skip List

- What if you add 'fast lanes'?
- Idea: 1 move in level  $k \approx 2$  in level  $k - 1 \approx 2^k$  in level 0
- $\mathcal{O}(\log n)$  expected search time
- Expected height of  $\mathcal{O}(\log n)$




Figure 1: A skip list with four coin flips

# Skip List

- What if you add 'fast lanes'?
- Idea: 1 move in level  $k \approx 2$  in level  $k - 1 \approx 2^k$  in level 0
- $\mathcal{O}(\log n)$  expected search time
- Expected height of  $\mathcal{O}(\log n)$
- Expected  $2n$  nodes




Figure 1: A skip list with four coin flips

# Skip List

- What if you add 'fast lanes'?
- Idea: 1 move in level  $k \approx 2$  in level  $k - 1 \approx 2^k$  in level 0
- $\mathcal{O}(\log n)$  expected search time
- Expected height of  $\mathcal{O}(\log n)$
- Expected  $2n$  nodes - bad





Figure 1: A skip list with four coin flips


# Zip Tree

- Idea: Construct a BST from a skip-list





# Zip Tree

- Idea: Construct a BST from a skip-list




# Zip Tree

- Idea: Construct a BST from a skip-list



# Zip Tree



- Idea: Construct a BST from a skip-list







# Zip Tree

- Idea: Construct a BST from a skip-list





# Zip Tree

- Idea: Construct a BST from a skip-list





# Zip Tree

- Idea: Construct a BST from a skip-list




# Zip Tree

- Idea: Construct a BST from a skip-list





# Zip Tree

- Idea: Construct a BST from a skip-list



# Zip Tree

- Idea: Construct a BST from a skip-list



# Zip Tree Analysis

- Max key expected depth  $\log n$
- Min key expected depth  $\log n/2$
- Asymmetric!

# Zip Tree Analysis

- Max key expected depth  $\log n$
- Min key expected depth  $\log n/2$
- Asymmetric!
- Results in average node depth of  $1.5 \log n$



# Zip Tree Analysis

- Max key expected depth  $\log n$
- Min key expected depth  $\log n/2$
- **Asymmetric!**
- Results in average node depth of  $1.5 \log n$
- Can we do better?

# Zip-zip Trees

- What if rank was a tuple,  $(r_1, r_2)$ ?

# Zip-zip Trees

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed
  - Let  $r_2$  be uniformly distributed from  $[1, \log^c n]$

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed
  - Let  $r_2$  be uniformly distributed from  $[1, \log^c n]$
- Compare ranks lexicographically

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed
  - Let  $r_2$  be uniformly distributed from  $[1, \log^c n]$
- Compare ranks lexicographically
- Metadata size  $O(\log \log n)$ ?

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed
  - Let  $r_2$  be uniformly distributed from  $[1, \log^c n]$
- Compare ranks lexicographically
- ✓ Metadata size  $O(\log \log n)$ ? -  $O(\log \log n) + O(c \log \log n)$

- What if rank was a tuple,  $(r_1, r_2)$ ?
  - Let  $r_1$  be geometrically distributed
  - Let  $r_2$  be uniformly distributed from  $[1, \log^c n]$
- Compare ranks lexicographically
- ✓ Metadata size  $O(\log \log n)$ ? -  $O(\log \log n) + O(c \log \log n)$
- Hope: fewer collisions, better depth?



# Zip-zip Trees Example




Figure 2: A zip tree

# Zip-zip Trees Example




Figure 2: A zip tree

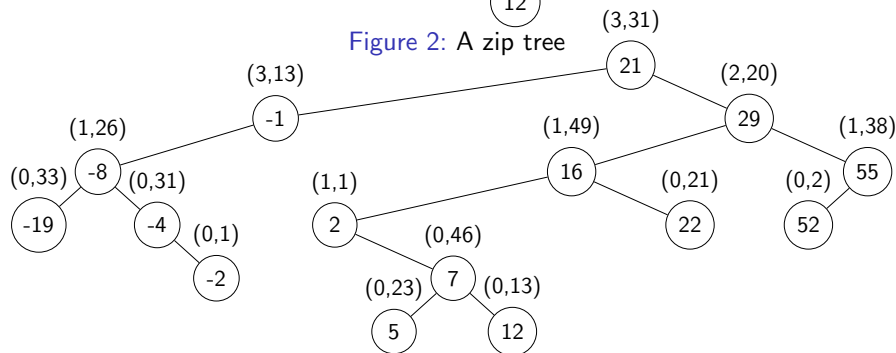




Figure 3: A random zip-zip tree generated from the above zip tree


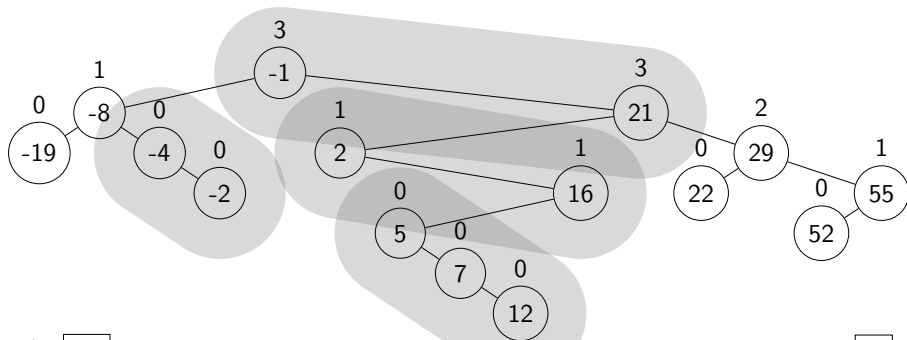
# Zip-zip Trees Analysis I

- Idea: Consider rank groups




# Zip-zip Trees Analysis I

- Idea: Consider rank groups





# Zip-zip Trees Analysis I

- Idea: Consider rank groups



# Zip-zip Trees Analysis I

- Idea: Consider rank groups



# Zip-zip Trees Analysis II


- How big are rank groups?

## Lemma

*The size of an  $r_1$  rank group has expected value 2 and is  $< 2 \log n$  w.h.p.*

## Proof Sketch.

- Size is (at most) geometrically distributed



# Zip-zip Trees Analysis III

## Theorem


*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*



# Zip-zip Trees Analysis III

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*




# Zip-zip Trees Analysis III

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$




# Zip-zip Trees Analysis III

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$




# Zip-zip Trees Analysis III

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$




# Zip-zip Trees Analysis III

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$



# Zip-zip Trees Analysis III


## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$

$$P(x_i \prec x_j) = \frac{1}{|i - j| + 1}$$



# Zip-zip Trees Analysis III

## Theorem


The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$

## Proof Sketch.

- A node  $x_i$  is an ancestor of node  $x_j$  iff  $x_i$  has maximum rank in  $[i, j]$

$$P(x_i \prec x_j) = \frac{1}{|i - j| + 1}$$

$$\delta_j = \sum_{i=1}^j \frac{1}{j - i + 1} + \sum_{i=j+1}^n \frac{1}{i - j + 1} = H_j + H_{n-j+1} - 1$$



# Zip-zip Trees Analysis IIII

## Theorem

*The expected depth,  $\delta_j$ , of the  $j$ -th smallest key in a zip-zip tree is  $H_j + H_{n-j+1} - 1 + o(1)$*

## Corollary

*The expected depth of the min and max keys is  $0.69 \log n + \gamma + o(1)$*

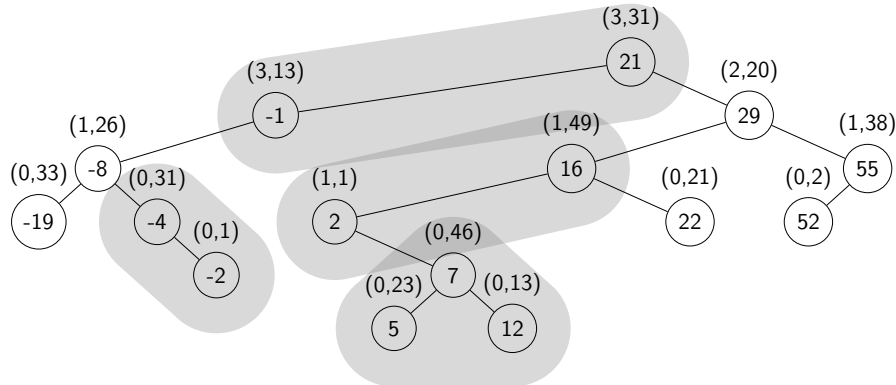
## Corollary

*The expected depth of any key is at most  $1.39 \log n - 1 + o(1)$*



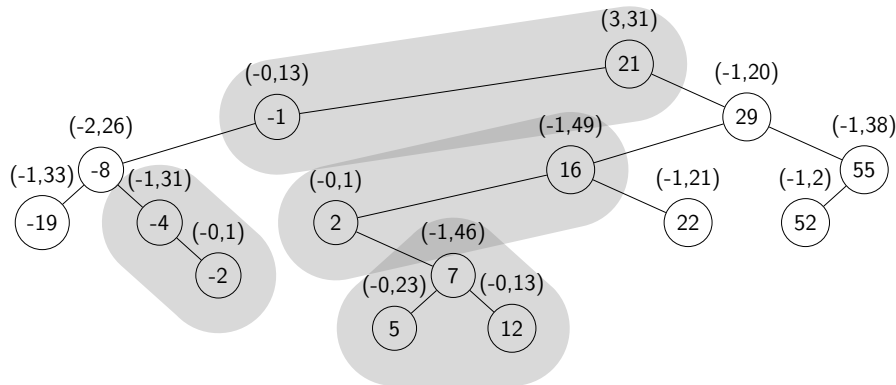
# Just-in-Time (JIT) Zip-zip Trees

- Ranks can be up to  $O(\log n)$ , but don't differ much




## Just-in-Time (JIT) Zip-zip Trees

- Ranks can be up to  $O(\log n)$ , but don't differ much
  - Store  $r_1$  rank differences! (Expected  $O(1)$ )



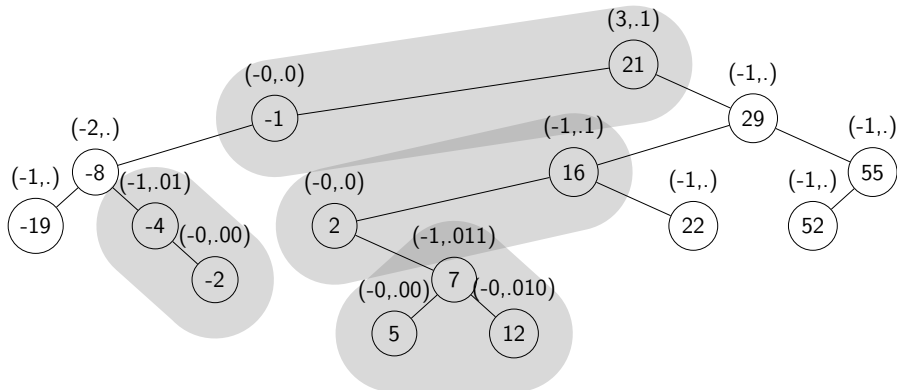
# Just-in-Time (JIT) Zip-zip Trees

- Ranks can be up to  $O(\log n)$ , but don't differ much
  - Store  $r_1$  rank differences! (Expected  $O(1)$ )
- Rank groups are small...



## Just-in-Time (JIT) Zip-zip Trees

- Ranks can be up to  $O(\log n)$ , but don't differ much
  - Store  $r_1$  rank differences! (Expected  $O(1)$ )
- Rank groups are small...
  - Generate  $r_2$  ranks on the fly! (Expected  $O(1)$ )<sup>1</sup>



<sup>1</sup> $r_1$  differences are  $O(1)$  per node,  $r_2$  are  $O(1)$  per operation

# Depth Discrepancy




Figure 4: The depth discrepancy between the min and max keys for three variants

# Average Key Depth and Tree Height





Figure 5: The average node depth and tree height for three variants

# Rank Collisions



**Figure 6:** The frequency of encountered rank ties per rank comparison for the uniform variant and per element insertion for the zip-zip variant

# Just-in-Time Zip-zip Tree Size




Figure 7: The metadata size for the just-in-time implementation



- 2023 – Zip-zip tree
  - ✓ 1.39  $\log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement
  - ✓ Strongly history independent (except JIT)
  - ✓ May be partially persistent
  - ✓ Supports biased keys (still  $O(\log \log n)$  bits per node)

# Summary

- 2023 – Zip-zip tree
  - ✓ 1.39  $\log n$  expected average depth
  - ✓ Space cost:  $\log \log n$  bits per node (or  $O(1)$  bits per update w.h.p.)
  - ✓ Easy to implement
  - ✓ Strongly history independent (except JIT)
  - ✓ May be partially persistent
  - ✓ Supports biased keys (still  $O(\log \log n)$  bits per node)
- Any questions?